

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")
```

```
from google.colab import files
uploaded = files.upload()
```

Churn_Modelling.csv

- **Churn_Modelling.csv**(text/csv) - 684858 bytes, last modified: 10/12/2022 - 100% done
Saving Churn_Modelling.csv to Churn_Modelling (1).csv

```
df=pd.read_csv("/content/Churn_Modelling.csv")
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname               10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender               10000 non-null  object
6   Age                  10000 non-null  int64
7   Tenure               10000 non-null  int64
8   Balance              10000 non-null  float64
9   NumOfProducts        10000 non-null  int64
10  HasCrCard            10000 non-null  int64
```

```

11 IsActiveMember    10000 non-null    int64
12 EstimatedSalary    10000 non-null    float64
13 Exited             10000 non-null    int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```

Univariate analysis

Histogram


```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.88928
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.40520
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.54000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.24000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.09000




```
sns.displot(df["Age"], color='darkorange')
```

```
<seaborn.axisgrid.FacetGrid at 0x7f36bbc6aa50>
```




```
sns.histplot(y="Age",data=df,color='darkorange')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36b89be150>
```



```
sns.histplot(x='Age',data=df,hue=df['Tenure'])
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36b8485950>
```



Distplot


```
sns.distplot(df["Age"],color='purple')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36b84ad910>
```



```
sns.distplot(df["Age"],hist=False,color='blue')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f36b8a9fe90>



Box plot

+ Code

+ Text

```
sns.boxplot(df["Age"],color='pink')
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36b7a51890>
```

Count plot



```
sns.countplot(df['Age'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36b79139d0>
```




Bivariate analysis

Bar plot

```
sns.barplot(df["NumOfProducts"],df["Age"])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36b7711050>
```




Lineplot



```
sns.lineplot(df["Age"],df["NumOfProducts"], color='purple')
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36b7689350>
```



Scatterplot

```
sns.scatterplot(x=df.Age,y=df.RowNumber,color='green')
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36b7aa9cd0>
```



Pointplot

```
sns.pointplot(x='Age',y='Tenure',data=df,color='pink')
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36b7bf8790>
```



Regplot

```
sns.regplot(df['Age'],df['Tenure'],color='orange')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f36b5bb7790>




Multi-variate analysis

Pairplot

```
sns.pairplot(data=df[["RowNumber", "Age", "Tenure", "Balance", "NumOfProducts"]],kind="kde")
```


<seaborn.axisgrid.PairGrid at 0x7f36b5aebfd0>



4.0
--


```
sns.pairplot(data=df[["RowNumber", "Age", "Tenure", "Balance", "NumOfProducts"]], hue="Age", diag_kind="hist")
```

<seaborn.axisgrid.PairGrid at 0x7f36b4c8de90>



```
sns.pairplot(data=df[["RowNumber", "Age", "Tenure", "Balance", "NumOfProducts"]], hue="Age")
```

<seaborn.axisgrid.PairGrid at 0x7f36b48bf190>



Descriptive statistics

df.describe()

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMem
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000



Handling the missing values



```
data=pd.DataFrame({"a":[1,2,np.nan],"b":[1,np.nan,np.nan],"c":[1,2,4]})
```

```
data
```

	a	b	c
0	1.0	1.0	1
1	2.0	NaN	2
2	NaN	NaN	4

```
data.isnull().any()
```

```
a      True
b      True
c     False
dtype: bool
```

```
data.isnull().sum()
```

```
a      1
b      2
c      0
dtype: int64
```

```
data.fillna(value = "S")
```

	a	b	c
0	1.0	1.0	1
1	2.0	S	2
2	S	S	4

```
data["a"].mean()
```

1.5

```
data["a"].median()
```

1.5

Finding and replacing outliers

```
outlierss=df.quantile(q=(0.25,0.75))
```


outlierss

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Ex
0.25	2500.75	15628528.25	584.0	32.0	3.0	0.00	1.0	0.0	0.0	51002.1100	
0.75	7500.25	15753233.75	718.0	44.0	7.0	127644.24	2.0	1.0	1.0	149388.2475	

```
sns.boxplot(df["Age"],color='purple')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36ae44cc50>
df["Age"]=np.where(df["Age"]<25,50,df["Age"])
sns.boxplot(df["Age"],color='pink')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36ae421e50>
```



Check for Categorical columns and perform encoding.

```
df.head(4)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	


```
df["Gender"].replace({"Female":0,"Male":1},inplace = True)
df["Geography"].replace({"France":1,"Spain":2,"Germany":3},inplace = True)
df["Gender"].replace({"Female":0,"Male":1},inplace = True)
df["Geography"].replace({"France":1,"Spain":2,"Germany":3},inplace = True)
```

```
df.head(4)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMei
0	1	15634602	Hargrave	619	1	0	42	2	0.00	1	1	
1	2	15647311	Hill	608	2	0	41	1	83807.86	1	0	
2	3	15619304	Onio	502	1	0	42	8	159660.80	3	1	
3	4	15701354	Boni	699	1	0	39	1	0.00	2	0	

Split the data into dependent and independent variables.

```
y = df["Surname"]
```

```
x=df.drop(columns=["Surname"],axis=1)
```

```
x.head()
```

RowNumber	CustomerId	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
-----------	------------	-------------	-----------	--------	-----	--------	---------	---------------	-----------	----------------	-----------------

Scale the independent variables

1	2	15647311	608	2	0	41	1	83807.86	1	0	1
---	---	----------	-----	---	---	----	---	----------	---	---	---

```
names=x.columns
```

```
names
```

```
Index(['RowNumber', 'CustomerId', 'CreditScore', 'Geography', 'Gender', 'Age',
      'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember',
      'EstimatedSalary', 'Exited'],
      dtype='object')
```

```
from sklearn.preprocessing import scale
```

```
X=scale(x)
```

```
X
```

```
array([[ -1.73187761, -0.78321342, -0.32622142, ...,  0.97024255,
         0.02188649,  1.97716468],
       [ -1.7315312 , -0.60653412, -0.44003595, ...,  0.97024255,
         0.21653375, -0.50577476],
       [ -1.73118479, -0.99588476, -1.53679418, ..., -1.03067011,
         0.2406869 ,  1.97716468],
       ...,
       [  1.73118479, -1.47928179,  0.60498839, ...,  0.97024255,
        -1.00864308,  1.97716468],
       [  1.7315312 , -0.11935577,  1.25683526, ..., -1.03067011,
        -0.12523071,  1.97716468],
       [  1.73187761, -0.87055909,  1.46377078, ..., -1.03067011,
        -1.07636976, -0.50577476]])
```

```
x = pd.DataFrame(X,columns = names )
```

```
x
```

	RowNumber	CustomerId	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive
0	-1.731878	-0.783213	-0.326221	-0.902587	-1.095988	0.179622	-1.041760	-1.225848	-0.911583	0.646092	0
1	-1.731531	-0.606534	-0.440036	0.301665	-1.095988	0.080092	-1.387538	0.117350	-0.911583	-1.547768	0
2	-1.731185	-0.995885	-1.536794	-0.902587	-1.095988	0.179622	1.032908	1.333053	2.527057	0.646092	-1
3	-1.730838	0.144767	0.501521	-0.902587	-1.095988	-0.118968	-1.387538	-1.225848	0.807737	-1.547768	-1
4	-1.730492	0.652659	2.063884	0.301665	-1.095988	0.279152	-1.041760	0.785728	-0.911583	0.646092	0
...
9995	1.730492	-1.177652	1.246488	-0.902587	0.912419	-0.118968	-0.004426	-1.225848	0.807737	0.646092	-1
9996	1.730838	-1.682806	-1.391939	-0.902587	0.912419	-0.517088	1.724464	-0.306379	-0.911583	0.646092	0
9997	1.731185	-1.479282	0.604988	-0.902587	-1.095988	-0.417558	0.687130	-1.225848	-0.911583	-1.547768	0
9998	1.731531	-0.119356	1.256835	1.505917	0.912419	0.179622	-0.695982	-0.022608	0.807737	0.646092	-1
9999	1.731878	-0.870559	1.463771	-0.902587	-1.095988	-1.213798	-0.350204	0.859965	-0.911583	0.646092	-1

10000 rows × 13 columns



Split the data into training and testing

```
from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)

x_train.head()
```

	RowNumber	CustomerId	CreditScore	Geography	Gender	Age	Tenure	B
7389	0.827747	-0.195066	0.170424	0.301665	-1.095988	-0.616618	-0.004426	-1.0
9275	1.481077	0.810821	-2.312802	1.505917	0.912419	0.179622	-1.387538	-0.0
2995	-0.694379	-1.507642	-1.195351	-0.902587	-1.095988	-1.114268	-1.041760	0.0
5316	0.109639	1.243462	0.035916	0.301665	0.912419	-0.019438	-0.004426	0.0
356	-1.608556	-1.100775	2.063884	0.301665	-1.095988	1.672571	1.032908	0.0

```
x_train.shape,y_train.shape,x_test.shape,y_test.shape
```

```
((8000, 13), (8000,), (2000, 13), (2000,))
```

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 11:56 AM

